

Lab 5. Anna Startseva BS19-04

I created the database "customers" using python script (warm_up.py). You can see it in repository.

The first query shows all ages that are repeated more than 2500 times. See its cost in "Seq Scan...":

1	EXPLAIN ANALYZE SELECT age FROM customer GROUP BY age HAVING count(age) > 2500;
<div>Data Output Explain Messages Notifications</div>	
	QUERY PLAN text
1	HashAggregate (cost=4590.00..4590.51 rows=14 width=4) (actual time=69.071..69.075 rows=6 loops=1)
2	Group Key: age
3	Filter: (count(age) > 2500)
4	Batches: 1 Memory Usage: 24kB
5	Rows Removed by Filter: 35
6	-> Seq Scan on customer (cost=0.00..4090.00 rows=100000 width=4) (actual time=0.019..16.306 rows=100000 loops=1)
7	Planning Time: 0.148 ms

Creating a b-tree index on the column "age":

index1 ✕

General

Definition

SQL

Name

index1

Tablespace

pg_default ✕ ▼

Comment

✕ Cancel 🔄 Reset 💾 Save

index1 ✕

General

Definition

SQL

Access Method

btree ▼

Fill factor

Unique?

☐ No

Clustered?

☐ No

Concurrent build?

☐ No

Constraint

1

Columns

+

Column	Operator class	Sort order	NULLs	Collation
age ▼	▼	<input checked="" type="checkbox"/> ASC	<input checked="" type="checkbox"/> LAST	▼

Include columns

age

✕ Cancel 🔄 Reset 💾 Save

The cost of the query after creating b-tree index is lower:

1	EXPLAIN ANALYZE SELECT age FROM customer GROUP BY age HAVING count(age) > 2500;
Data Output Explain Messages Notifications	
QUERY PLAN text	
1	GroupAggregate (cost=0.42..3112.93 rows=14 width=4) (actual time=5.583..41.017 rows=6 loops=1)
2	Group Key: age
3	Filter: (count(age) > 2500)
4	Rows Removed by Filter: 35
5	-> Index Only Scan using index1 on customer (cost=0.42..2612.42 rows=100000 width=4) (actual time=0.335..24.297 rows=100000 l...
6	Heap Fetches: 0
7	Planning Time: 6.261 ms

The second query shows all names of the customers which contain letter 'W'. See its cost in "Seq Scan...":

1	EXPLAIN ANALYZE SELECT name FROM customer WHERE name LIKE '%W%';
Data Output Explain Messages Notifications	
QUERY PLAN text	
1	Seq Scan on customer (cost=0.00..4340.00 rows=6232 width=14) (actual time=0.015..25.938 rows=9681 loops=1)
2	Filter: (name ~~ '%W%':text)
3	Rows Removed by Filter: 90319
4	Planning Time: 0.107 ms
5	Execution Time: 26.356 ms

Creating a hash index on the column "name":

index2 ✕

General

Definition

SQL

Name

index2

Tablespace

pg_default ✕ ▼

Comment

✕ Cancel 🔄 Reset 💾 Save

index2 ✕

General

Definition

SQL

Access Method

hash ▼

Fill factor

Unique?

☐ No

Clustered?

☐ No

Concurrent build?

☐ No

Constraint

1

Columns

+

Column	Operator class	Sort order	NULLs	Collation
name ▼	▼	<input checked="" type="checkbox"/> ASC	<input type="checkbox"/> LAST	pg_catalog."default" ▼

Include columns

✕ Cancel 🔄 Reset 💾 Save

The cost of the query after creating index doesn't change because hash indexes only effective when there is '=' operation:

```
1 EXPLAIN ANALYZE SELECT name FROM customer WHERE name LIKE '%W%';
```

Data Output Explain Messages Notifications

QUERY PLAN
text

1	Seq Scan on customer (cost=0.00..4340.00 rows=6232 width=14) (actual time=0.029..36.514 rows=9681 loops=1)
2	Filter: (name ~~ '%W%':text)
3	Rows Removed by Filter: 90319
4	Planning Time: 7.727 ms
5	Execution Time: 37.048 ms

So, I created the query with '=' operation that finds all customers with name "Anna Watson". See its cost in "Seq Scan...":

```
1 EXPLAIN ANALYZE SELECT * FROM customer WHERE name = 'Anna Watson';
```

Data Output Explain Messages Notifications

	QUERY PLAN
1	Seq Scan on customer (cost=0.00..4340.00 rows=2 width=215) (actual time=9.111..15.911 rows=2 loops=1)
2	Filter: (name = 'Anna Watson':text)
3	Rows Removed by Filter: 99998
4	Planning Time: 0.121 ms
5	Execution Time: 15.929 ms

The cost of the query after creating hash index is lower:

1	EXPLAIN ANALYZE SELECT * FROM customer WHERE name = 'Anna Watson';
<div>Data Output Explain Messages Notifications</div>	
	<div>QUERY PLAN</div> <div>text</div>
1	Bitmap Heap Scan on customer (cost=4.02..11.89 rows=2 width=215) (actual time=0.041..0.044 rows=2 loops=1)
2	Recheck Cond: (name = 'Anna Watson'::text)
3	Heap Blocks: exact=2
4	-> Bitmap Index Scan on index2 (cost=0.00..4.01 rows=2 width=0) (actual time=0.033..0.033 rows=2 loops=1)
5	Index Cond: (name = 'Anna Watson'::text)
6	Planning Time: 6.015 ms
7	Execution Time: 0.075 ms